

# 复旦大学软件学院

## 2010~2011 学年第二学期期末考试试卷

### Problem 1: (16 points)

1. Consider the following C code and assembly code for a strange but simple function on x86(32-bit machine) platform:

<pre> int lol(int a, int b) {     switch(a){         case 210:             b *= 13;             ____[1]____;         case 213:             b = 18243;         case 214:             b *= b;             ____[2]____;         case 216:         case 218:             b -= a;             ____[3]____;         case 219:             b += 17;             ____[4]____;         default:             b = -9;     }     return b; } </pre>	<pre> &lt;lol&gt;: 0x8048330: push    %ebp 0x8048331: mov     %esp,%ebp 0x8048333: mov     0x8(%ebp),%eax 0x8048336: mov     0xc(%ebp),%ecx 0x8048339: lea     0xffffffff2e(%eax),%edx 0x804833f: cmp     \$0x8,%edx 0x8048342: jbe     0x8048350 0x8048345: mov     \$0xfffffffff7,%eax 0x804834a: ret 0x8048350: jmp     *0x8048470(,%edx,4) 0x8048358: sub     %eax,%ecx 0x804835a: mov     %ecx,%eax 0x8048360: ret 0x8048361: mov     %ecx,%eax 0x8048363: imul    %ecx,%eax 0x8048367: ret 0x8048369: mov     \$0x13d63b89,%eax 0x804836e: ret 0x8048370: lea     (%ecx,%ecx,2),%eax 0x8048373: lea     (%ecx,%eax,4),%eax 0x8048376: ret </pre>
-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

- For each case in the switch block which should have a **"break"**, write it on the corresponding blank line. ( $2' * 4 = 8'$ )  
(**Hint:**  $0xffffffff2e = -210$  and  $0x13d63b89 = 18243 * 18243$ )
- Using the available information, fill in the jump table below. ( $2' * 4 = 8'$ )

Jump table entries:	
0x8048470:	____[5]____
0x8048474:	0x08048345
0x8048478:	0x08048345
0x804847c:	____[6]____
0x8048480:	____[7]____
0x8048484:	0x08048345
0x8048488:	____[8]____
0x804848c:	0x08048345
0x8048490:	0x08048358
0x8048494:	0x00000000

## Problem 2: (27points)

Suppose the following C code and assembly code are executed on a 32-bit **little endian** machine. Please read the following code carefully and answer questions.

### Assembly Code:

<b>main:</b>		
3d:	8d 4c 24 04	lea 0x4(%esp),%ecx
41:	83 e4 f0	and \$0xffffffff0,%esp
44:	ff 71 fc	pushl -0x4(%ecx)
47:	55	push %ebp
48:	89 e5	mov %esp,%ebp
4a:	51	push %ecx
4b:	83 ec 28	sub \$0x28,%esp
4e:	c6 45 ec 12	movb \$0x12,-0x14(%ebp)
52:	c6 45 ed 34	movb \$0x34,-0x13(%ebp)
56:	c6 45 ee 56	movb \$0x56,-0x12(%ebp)
5a:	c6 45 ef 78	movb \$0x78,-0x11(%ebp)
5e:	66 c7 45 e8 ____ [1] ____	movw \$0x1234,-0x18(%ebp)
64:	66 c7 45 ea ____ [2] ____	movw \$0x5678,-0x16(%ebp)
6a:	c7 45 f0 ____ [3] ____	movl \$0x78563412,-0x10(%ebp)
71:	c7 45 f4 ____ [4] ____	movl \$0x12345678,-0xc(%ebp)
78:	8d 45 ec	lea ____ [5] ____,%eax
7b:	89 45 f8	mov %eax,-0x8(%ebp)
7e:	8d 45 e8	lea -0x18(%ebp),%eax
81:	89 44 24 04	mov %eax,0x4(%esp)
85:	8b 45 f8	mov -0x8(%ebp),%eax
88:	89 04 24	mov %eax,(%esp)
8b:	e8 fc ff ff ff	call 8c <main+0x4f>
90:	89 c2	mov %eax,%edx
92:	8b 45 f8	mov -0x8(%ebp),%eax
95:	89 10	mov %edx,(%eax)
97:	83 c4 28	add \$0x28,%esp
9a:	59	pop %ecx
9b:	5d	pop %ebp
9c:	8d 61 fc	lea -0x4(%ecx),%esp
9f:	c3	ret

### C Code:

```
int func(int *addr, int* save)
{
    int x = *addr;
    int y = *(addr + 1);

    y = x & y;
    *save = y;
    x = *(addr + 2);
    y = x | y;

    return y;
}

int main(void)
{
    char a[] = {0x12,0x34,0x56,0x78};
    short b[] = {0x1234,0x5678};
    int c = 0x78563412;
    int d = 0x12345678;
    int *e = (int*)a;

    *e = func(e, (int*)b);
}
```


1. Please fill the blanks in the assembly code according to C Code with the correct **byte ordering**. (1' \* 5 = 5')
2. According to the C code, please fill the following blanks with "<", "=" or ">". (2' \* 3 = 6')

\*((int\*)a) \_\_\_\_ [1] \_\_\_\_ \*((int\*)b)

\*((int\*)a) \_\_\_\_ [2] \_\_\_\_ c

\*((int\*)a) \_\_\_\_ [3] \_\_\_\_ d

3. Just before the call instruction in main function has been executed, the **%esp** value is **0xbfaa42c0**. Please fill the following stack diagram with correct value in hex format. (NOTE: If the value is undetermined, fill "-" instead) (2' \* 4 = 8')



4. Suppose the statement "**\*e = func(e, (int \*)b);**" in main() function has been executed, please fill the following table. (2' \* 4 = 8')

a[0]	a[2]	b[0]	b[1]
__[1]__	__[2]__	__[3]__	__[4]__

### Problem 3: (27 points)

1. A software school student decided to write a dynamic memory allocator for an **x86-64** machine in which each block has the following form:

Header	Id String	<b>Payload</b>	Footer
--------	-----------	----------------	--------

Where

- Header is a **4** byte header
- Id string is an **8** byte string
- Payload is arbitrary size, including padding
- Footer is a **4** byte footer

Assume the student wants to print the Id string with the following function

- 1) Where bp points to the beginning of the **Payload** and is aligned to 0 modulo 8. Write **"YES"** to each of the following 4 macros if it will correctly print the Id string, else write **"NO"** ( $2' * 4 = 8'$ )

a) #define GET\_ID(bp) ((char \*)(((long)bp)-8))          **[1]**      
b) #define GET\_ID(bp) ((char \*)(((int)bp)-8))          **[2]**      
c) #define GET\_ID(bp) ((char \*)(((long \*)bp)-1))          **[3]**      
d) #define GET\_ID(bp) ((char \*)(((int \*\*)bp)-2))          **[4]**    

- 2) The student uses **Header** to encode the block size (including all padding). If we impose a **one-word** alignment (8 bytes) constraint, then how many bits of **Header** can be used to indicate whether the block is allocated or free? ( $2'$ )

2. Assume we use the **buddy** system to manage memory. In the buddy system, a block of memory size  $2^i$ , is called an **i-block**, and the **i-list** consists of starting addresses of **free** i-blocks. Suppose that the total memory size in our computer system is **1664** bytes, whose starting address is **0**. ( $18'$ )

- 1) Originally, all these 1664 bytes are free, and are represented as **3** free blocks, and sizes are 1024 bytes, 512 bytes and 128 bytes. What are the originally 3 free blocks? ( $2' * 3 = 6'$ )

    **[1]**    -block ,     **[2]**    -block ,     **[3]**    -block

- 2) Suppose we have to issue a request for allocating 100 bytes, then we should check     **[4]**    -list first, and find there are/is     **[5]**     free block(s) in that list. ( $2' * 2 = 4'$ )

3) If your answer to **[5]** is 0, then we will split the free block in **\_\_[6]\_\_**-list to satisfy the 100 bytes allocating request; **else**, we don't need to split any free blocks, and please write NULL to **[6]**. (2')

4) After issue the 100 bytes request, suppose the following memory requests are issued sequentially:

1. malloc(128)
2. malloc(28)
3. malloc(56)
4. malloc(110)
5. malloc(60)
6. malloc(48)
7. malloc(230)

Tell the number of free blocks in each of the free lists in the following table, after these memory requests are processed. (1' \* 6 = 6')

Free list	Number of free blocks in the list
5-list	[7]
6-list	[8]
7-list	[9]
8-list	[10]
9-list	[11]
10-list	[12]

## Problem 4: (30 points)

1. Please optimize the following function **func()**. The function **length(x)** returns the length of vector **x**, and the function **get\_element(x,y)** returns the **y**th element of vector **x**. (15')

The function **func()** is to calculate:

$$*dest += \sum_{i=1} (a[i-1] + a[i-1] + a[i-1] + a[i-1]) * (b[i-1] + b[i-1] + b[i-1] + b[i-1])$$

```
int get_element(int *x, y)
{
    if(y<1 || y> length(x))
        return -1;

    int ret= x[y-1];
    return ret;
}

int func(int* a, int* b, int *dest)
{
    int i,j,k;
    if(length(a) != length(b))
        return 0;

    for (i = 1; i <= length(a); i++) {
        int tempMul, val1_mul2 = 0, val2_mul2 = 0;
        int tempSum = get_element(a,i) + get_element(b,i);

        for (j=0;j<4;++j) {
            val1_mul2 = val1_mul2+get_element(a, i);
        }
        for(k=0;k<4;++k) {
            val2_mul2 = val2_mul2+get_element(b, i);
        }

        tempMul= val1_mul2 * val2_mul2;
        *dest = *dest+ tempMul;
    }

    return 1;
}
```

2. Consider the following function. Assume that there is a single cache with an **8** bytes block size. "A" is a N\*M array of **int** (the size of int is 4 bytes). **The array size M and L are so large that a single matrix row does not fit in the cache.** All loop indexes (i, j and k) are stored in registers, which not require any load or store instructions. The memory alignment and compiler-driven optimizations are **\_NOT\_** considered.

```

#define N 10
#define M 20
#define L 20

int A[N][M];
int B[L];

int func(void)
{
    int ret=0, i, j, k;
    for (i = 0 ; i < N ; i ++){
        for (j = 0 ; j < M ; j ++){
            for (k = 0 ; k < L ; k ++){
                if (A[i][j] + B[k] > 0 )
                    ret ++;
            }
        }
    }
    return ret;
}

```

- 1) What is the expected cache miss rate of the access to A and B? (6')  
A: \_\_\_\_ [1] \_\_\_\_                      B: \_\_\_\_ [2] \_\_\_\_
- 2) Fill in the blanks of code below with the most cache friendly iteration order you think. (e.g. i, j, k, N, M, L) (3')

```

#define N 10
#define M 20
#define L 20

int A[N][M];
int B[L];

int func2(void)
{
    int ret=0, i, j, k;
    for (____ [1] ____ = 0; ____ [1] ____ < ____ [2] ____; ____ [1] ____ ++){
        for (____ [3] ____ = 0 ; ____ [3] ____ < ____ [4] ____; ____ [3] ____ ++){
            for (____ [5] ____ = 0 ; ____ [5] ____ < ____ [6] ____; ____ [5] ____ ++){
                if (A[i][j] + B[k] > 0 )
                    ret ++;
            }
        }
    }
    return ret;
}

```

- 3) What is the expected cache miss rate of the access to A and B? (6')  
A: \_\_\_\_ [1] \_\_\_\_                      B: \_\_\_\_ [2] \_\_\_\_



## Solution

### Problem 1: (16points)

- |    |     |     |
|----|-----|-----|
| 1. | [1] | [2] |
|    | [3] | [4] |
|    | [5] | [6] |
|    | [7] | [8] |

### Problem 2: (27points)

- |   |     |     |     |
|---|-----|-----|-----|
| 1 | [1] | [2] |     |
|   | [3] | [4] |     |
|   | [5] |     |     |
| 2 | [1] | [2] | [3] |
| 3 | [1] | [2] |     |
|   | [3] | [4] |     |
| 4 | [1] | [2] |     |
|   | [3] | [4] |     |

### Problem 3: (27points)

- |   |        |      |      |     |
|---|--------|------|------|-----|
| 1 | 1) [1] | [2]  | [3]  | [4] |
|   | 2)     |      |      |     |
| 2 | [1]    | [2]  | [3]  |     |
|   | [4]    | [5]  |      |     |
|   | [6]    |      |      |     |
|   | [7]    | [8]  | [9]  |     |
|   | [10]   | [11] | [12] |     |

### Problem 4: (30points)

1.

2.

1)      [1]                          [2]

2)      [1]                          [2]

[3] [4]

[5] [6]

3) [1] [2]