

上 海 交 通 大 学 试 卷 (A 卷)

(2011 至 2012 学 年 第 2 学 期)

班级号_____ 学号_____ 姓名 _____

课程名称 _____ 计算机系统基础 (1) _____ 成绩 _____

Problem 1: Y86 (13points)

1. [1] [2] [3]
[4] [5] [6]
2. [1] [2] [3] [4]
[5] [6] [7]

Problem 2: Memory (12points)

1 1)

2)

3)

Problem 3: x86-64 (20points)

- 1 [1] [2] [3] [4]
[5] [6] [7] [8]
[9] [10]

我承诺，我将严格遵守考试纪律。

承诺人：_____

题号	1	2	3	4	5				
得分									
批阅人(流水阅卷教师签名处)									

Problem 4: Cache (27points)

1. [1] [2] [3] [4]
2. [1] [2] [3] [4]
- [5] [6] [7] [8]
- [9] [10] [11] [12]
- [13]
3. 1)
- 2)
- 3)

Problem 5: Linking (28points)

1. [1] [2] [3]
- [4] [5]
2. [1] [2] [3]
- [4] [5] [6]
3. [1]
- [2]
- [3]
4. 1)
- 2)

Problem 1: Y86 (13 points)


0x000:	# Execution begins at address 0
0x000: 30f400010000	.pos 0
[1]: 30f500010000	init: irmovl Stack, %esp
0x00c: 7020000000	irmovl Stack, %ebp
	jmp Main
0x014:	# Array of 4 elements
0x014:	.align 4
0x014: 000b0000	array:
0x018: [2]	.long 0x00000b00
0x01c: [3]	.long 0x0fff6320
0x020: 30f214000000	.long 0x12345678
0x026: a02f	Main: [4]
0x028: 802e000000	pushl %edx
0x02d: 00	call Foo
0x02e: a05f	halt
0x030: 2045	Foo: pushl %ebp
0x032: 501508000000	rrmovl %esp, %ebp
0x038: 30f203000000	mrmovl 8(%ebp), %ecx
0x03e: 6222	irmovl \$3, %edx
0x040: [5]	andl %edx, %edx
0x045: 506100000000	je End
0x04b: 2060	Loop: mrmovl (%ecx), %esi
0x04d: 30f304000000	rrmovl %esi, %eax
0x053: 6031	irmovl \$4, %ebx
0x055: 30f3ffffffffff	addl %ebx, %ecx
0x05b: 6032	irmovl \$-1, %ebx
0x05d: 7445000000	addl %ebx, %edx
0x062: b05f	jne Loop
0x064: 90	End: [6]
0x100:	ret
0x100:	.pos 0x100
	Stack:

- 1) Fill in the blanks according to the Y86 assembly and object code. (1' * 6 = 6')
- 2) Fill in the value of the registers after the above program has **halted**. (1' * 7 = 7')

%eax: [1] %ebx: [2] %ecx: [3] %edx: [4]
 %esi: [5] %edi: -- %ebp: [6] %esp: [7]

Problem 2: Memory (12 points)

- The figure simulates the initial status of memory at a certain time. Allocated blocks are shaded, and free blocks are blank (each block represents **1 word = 4bytes**). Headers and footers are labeled with the number of bytes and allocated bit. The allocator maintains **double-word** alignment. Given the execution sequence of memory allocation operations (**malloc** or **free**) from 1 to 7. Please answer the following questions.



```

1:  P3 = malloc(16);
2:  P4 = malloc(8);
3:  P5 = malloc(3);
4:  free(P3);
5:  free(P4);
6:  free(P2);
7:  P6 = malloc(24);

```

- Assume **first-fit** algorithm is used to find free blocks. Please draw the status of memory and mark with variables after the **3rd** operation is executed. (4')
- Assume **best-fit** algorithm is used to find free blocks. Please draw the status of memory and mark with variables after the **3rd** operation is executed. (4')
- To analyze the memory utilization and time utilization of **first-fit** and **best-fit** in above execution sequence. **Which** is better and **why**? (4')

Problem 3: x86-64 (20 points)

Suppose the following C and assembly code is defined on a **64-bit little endian** machine (x86-64/Linux).

```
void foo (long int *pt1, long int *pt2, char *pt3, char *pt4,
          int in5, int in6, long int *pt7, long int *pt8)
{
    long int temp = *pt7 + *pt8;
}

int main(void)
{
    long int ia=1, ib=0xabcd;
    char ic=2, id=0x12;
    int ie=6, ig=7;
    long int ih=8, il=9;

    foo(&ia,&ib,&ic,&id,ie,ig,&ih,&il);

    return 0;
}
```

0000000000000034 <main>:

34:	55	push	%rbp
35:	48 89 e5	mov	%rsp,%rbp
38:	53	push	%rbx
39:	48 83 ec 40	sub	\$0x40,%rsp
3d:	48 c7 45 e8 01 00 00 00	movq	\$0x1,-0x18(%rbp)
45:	48 c7 45 e0 cd ab 00 00	movq	\$0xabcd,-0x20(%rbp)
4d:	c6 45 df 02	movb	\$0x2,-0x21(%rbp)
51:	c6 45 de 12	movb	\$0x12,-0x22(%rbp)
55:	c7 45 f0 06 00 00 00	movl	\$0x6,-0x10(%rbp)
5c:	c7 45 f4 07 00 00 00	movl	\$0x7,-0xc(%rbp)
63:	48 c7 45 d0 08 00 00 00	movq	\$0x8,-0x30(%rbp)
6b:	48 c7 45 c8 09 00 00 00	movq	\$0x9,-0x38(%rbp)
73:	44 8b 45 f4	mov	-0xc(%rbp),%r8d
77:	8b 7d f0	mov	-0x10(%rbp),%edi
7a:	48 8d 4d de	lea	-0x22(%rbp),__[1]__
7e:	48 8d 55 df	lea	-0x21(%rbp),__[2]__
82:	48 8d 5d e0	lea	-0x20(%rbp),%rbx
86:	48 8d 45 e8	lea	-0x18(%rbp),%rax
8a:	48 8d 75 c8	lea	__[3]__,%rsi
8e:	48 89 74 24 08	mov	%rsi,0x8(%rsp)
93:	48 8d 75 d0	lea	-0x30(%rbp),%rsi
97:	48 89 34 24	mov	%rsi,__[4]__
9b:	45 89 c1	mov	%r8d,%r9d
9e:	41 89 f8	mov	%edi,__[5]__

a1:	48 89 de	mov	%rbx,%rsi
a4:	48 89 c7	mov	%rax,__[6]__
a7:	e8 00 00 00 00	callq	foo <main+0x78>

0000000000000000 <foo>:			
0:	55	push	%rbp
1:	48 89 e5	mov	%rsp,%rbp
4:	48 89 7d e8	mov	%rdi,-0x18(%rbp)
8:	48 89 75 e0	mov	__[7]__, -0x20(%rbp)
c:	48 89 55 d8	mov	%rdx,-0x28(%rbp)
10:	48 89 4d d0	mov	__[8]__, -0x30(%rbp)
14:	44 89 45 cc	mov	%r8d,-0x34(%rbp)
18:	44 89 4d c8	mov	__[9]__, -0x38(%rbp)
1c:	48 8b 45 10	mov	0x10(%rbp),%rax
20:	48 8b 10	mov	(%rax),%rdx
23:	48 8b 45 18	mov	__[10]__,%rax
27:	48 8b 00	mov	(%rax),%rax
2a:	48 8d 04 02	lea	(%rdx,%rax,1),%rax
2e:	48 89 45 f8	mov	%rax,-0x8(%rbp)

1) Fill in the blanks according to the assembly code and c code. (2'*10 = 20')_

Problem 4: Cache (27 points)

Consider a **32-bit** machine with a **4-way** set associative cache. There are **16 sets**. Each block is **32 bytes**. Answer the following questions:

1. please fill the following blanks (4')

Cache size:__[1]__ bytes

Field	Length(bit)
Total	32
Tag	[2]
Set	[3]
Offset	[4]

2. Assume the cache is initially **empty**. There are several memory accesses in order. Fill the second table and compute the miss rate. ($1' * 12 + 2' = 14'$)

Order	Address	Set	Miss or not (Yes/No)
1	0xe166	[1]	[2]
2	0xe17e	[3]	[4]
3	0xe15a	[5]	[6]
4	0x9d60	[7]	[8]
5	0x9d5f	[9]	[10]
6	0xe142	[11]	[12]

Miss rate: [13]

3. Assume the following code is executed on this machine. You should also assume the following:

- ✧ `sizeof(int) = 4`
- ✧ the starting address of array `a` is `0x100`
- ✧ variable `i, j, t` are in registers

```
#define M 4
#define N 8

int a[M][N];
int i, j;

for (i = 0; i < M - 1; i++) {
    int t = a[i + 1][0];
    for (j = 0; j < N; j++)
        a[i][j] = a[i][j] * t;
}
```

- 1) Total number of memory accesses: [1] (2')
- 2) Miss rate: [2] (2')
- 3) Suppose you have the chance to **double** the cache size. **Which** following scheme is best choice for above code and **why?** (5')

Scheme A: double the number of sets. (i.e., 32 sets)

Scheme B: double the block size (i.e., 64 bytes)

Scheme C: double the number of ways (i.e., 8-way)

Problem 5: Linking (28 points)

The following program consists of three source files: node.h, display.c and main.c. The corresponding relocatable object files are also listed.

display.c	display.o
<pre>#include <stdio.h> #include "node.h" extern node_t head; extern node_t second; void display(void) { printf("%d %d", head.value, second.value); }</pre>	<pre>.text: 00000000 <display>: 0: 55 push %ebp 1: 89 e5 mov %esp,%ebp 3: 83 ec 18 sub \$0x18,%esp 6: a1 00 00 00 00 mov 0x0,%eax b: 8b 15 00 00 00 00 mov 0x0,%edx 11: 89 44 24 08 mov %eax,0x8(%esp) 15: 89 54 24 04 mov %edx,0x4(%esp) 19: c7 04 24 00 00 00 00 movl \$0x0, (%esp) 20: e8 fc ff ff ff call 21<display+0x21> 25: c9 leave 26: c3 ret</pre>
node.h and main.c	main.o
<pre>node.h: typedef struct node { int value; struct node* next; } node_t; ----- #include "node.h" extern void display(); node_t head, second; int main(void) { head.next=&second; head.value=1; head.next->value=2; second.next=0; display(); return 0; }</pre>	<pre>00000000 <main>: 0: 8d 4c 24 04 lea 0x4(%esp),%ecx 4: 83 e4 f0 and \$0xffffffff0,%esp 7: ff 71 fc pushl 0xffffffffc(%ecx) a: 55 push %ebp b: 89 e5 mov %esp,%ebp d: 51 push %ecx e: 83 ec 04 sub \$0x4,%esp 11: c7 05 04 00 00 00 00 movl \$0x0,0x4 18: 00 00 00 1b: c7 05 00 00 00 00 01 movl \$0x1,0x0 22: 00 00 00 25: a1 04 00 00 00 mov 0x4,%eax 2a: c7 00 02 00 00 00 movl \$0x2, (%eax) 30: c7 05 04 00 00 00 00 movl \$0x0,0x4 37: 00 00 00 3a: e8 fc ff ff ff call 3b <main+0x3b> 3f: b8 00 00 00 00 mov \$0x0,%eax 44: 83 c4 04 add \$0x4,%esp 47: 59 pop %ecx 48: 5d pop %ebp 49: 8d 61 fc lea 0xffffffffc(%ecx),%esp 4c: c3 ret</pre>

Partial `.symbol` table after relocation

Num	Value	Size	Type	Bind	Vis	Ndx	Name
52:	08049670	8	OBJECT	GLOBAL	DEFAULT	24	second
61:	080483a4	39	FUNC	GLOBAL	DEFAULT	12	display
68:	08049678	8	OBJECT	GLOBAL	DEFAULT	24	head
71:	080483cc	77	FUNC	GLOBAL	DEFAULT	12	main

1. The global variable "head" is referenced **three times** and "second" is referenced **twice** in `main.c` file. Please fill in the following relocation entries for symbol "head" and "second" in `main.o` ($1' * 5 = 5'$)

OFFSET	TYPE	SYMBOL
[1]	R_386_32	head
[2]	R_386_32	head
[3]	R_386_32	head
[4]	R_386_32	second
[5]	R_386_32	second

2. Please fill in the relocation entries in `display.o` ($2' * 6 = 12'$)

OFFSET	TYPE	SYMBOL
[1]	[2]	head
[3]	[4]	second
0x0000001c	R_386_32	.rodata
[5]	[6]	printf

3. After relocation and the program is built, what changes will happen to the underlined 3 instructions according to `.symbol` table given above? ($2' * 3 = 6'$)

b: 8b 15 00 00 00 00 (display.o mov 0x0,%edx)

After relocation: [1]

3a: e8 fc ff ff ff (main.o: call 3b <main+0x3b>)

After relocation: [2]

11: c7 05 04 00 00 00 00 00 00 00 (main.o: movl \$0x0,0x4)

After relocation: [3]

4. After relocation and the program is built, part of its procedure linkage table (PLT) is shown as follow:

080482b8 <printf@plt>:

80482b8:	ff 25 60 96 04 08	jmp	*0x8049660
80482be:	68 10 00 00 00	push	\$0x10
80482c3:	e9 c0 ff ff ff	jmp	8048288 <_init+0x18>

- 1) What is the **value** of word starting from address 0x08049660? (2')
- 2) Briefly describe the procedure of lazy binding when "printf" is called. (3')