

Introduction to Computer Systems

2015 Fall Middle Examination

Name_____ Student No._____ Score_____

Problem 1: (14 points)

[1]	[2]	[3]
[4]	[5]	[6]
[7]		

Problem 2: (10 points)

[1]	[2]
[3]	[4]
[5]	[6]
[7]	[8]
[9]	[10]

Problem 3: (16 points)

1. [1]	[2]
[3]	[4]
[5]	[6]
[7]	[8]
[9]	[10]
[11]	[12]

2.

3.

Problem 4: (10 points)

1 [1]	[2]
[3]	[4]

2

3.

Problem 5: (23 points)

1	[1]	[2]
	[3]	[4]
	[5]	[6]
	[7]	[8]
	[9]	[10]
	[11]	[12]
	[13]	[14]

2

Problem 6: (27 points)

1	[1]	[2]
	[3]	[4]
	[5]	[6]

2

3	[1]	[2]
	[3]	[4]
	[5]	[6]

4

5	[1]	[2]
	[3]	[4]

Problem 1: (14 points)

1. Consider the following C program

```
unsigned int ua = 0xe5;
short a = ua;
int b = a;
short c = (~a & 7) + 1;
int d = 0x69 && 0x55;
int e = ua >> 3;
```

Assume we are running code on an **8-bit** machine using two's complement arithmetic for signed integers. Also assume that right shifts of signed values are performed **arithmetically**. A "short" integer is encoded using **4 bits**. Please fill in the blanks of table below. ($2^8=16$)

Expression	Binary Representation
ua	1110 0101
a	[1]
b	[2]
c	[3]
e	[4]
d & (!0x41)	[5]
(e & 0x69) (a >> 2)	[6]
(e >> c) - 1 + (d << 1)	[7]

Problem 2: (10points)

Suppose a **32-bit little endian** machine has the following memory and register status. (NOTE: **Instructions are independent**). ($1^*10=10$)

Memory status

Address	Low	High
0x610	0xff 0x33 0x22 0x10	
0x614	0xab 0x00 0xff 0x11	
0x618	0x13 0x01 0xab 0xcd	
0x61c	0x01 0x54 0x76 0x98	

Register status

Register	Value
%eax	0x00000610
%ecx	0x00000001
%edx	0x00000003

Fill in the blanks using **4 byte size** and **hex**.

Operation	Destination	Value
addl %ecx, (%eax)	[1]	[2]
subl %edx, 4(%eax)	[3]	[4]
imul \$16, (%eax,%edx,4)	[5]	[6]
incl 8(%eax)	[7]	[8]
leal (%eax, %ecx,c), %eax	[9]	[10]

Problem 3: (16points)

Please answer the following questions according to the definition of the union. (NOTE that the size of different types in x86 and x86-64 is shown in the Figure 3.34 in ICS book.)

```
struct {
    short s1;
    short *ps[2];
    char ca[2];
    union {
        char c1;
        int *pi[2];
    } u;
    short s2;
    int (*p[2])();
    char c2;
} str[2];
```


1. Fill in the following blocks. (12')

Representation	x86	x86-64
sizeof(str[0])	[1]	[2]
sizeof(str[1].u)	[3]	[4]
str	0x8049780	0x600a00
&(str[0].u)	[5]	[6]
&(str[0].ca[1])	[7]	[8]
&(str[0].u.pi[1])	[9]	[10]
&(str[1].c2)	[11]	[12]

2. How many bytes are **wasted** in `str[0]` in x86 and x86-64? (1'*2=2')
3. If you can rearrange the declarations in the **struct** and **union**, how many bytes of memory can you **save** in `str[0]` compared to the original declaration in x86 and x86-64? (1'*2=2')

Problem 4: FP (10points)

The following figure shows the floating-point format we designed for the exam, called float12. Except for the length, it's the same as the IEEE 754 single-precision format you have learned in the class.



1. Fill the blanks with proper values. (4')

- 1) **Normalized:** $(-1)^S \times (1.\text{Fract}) \times 2^{\text{Exp.-bias}}$, where **bias**= [1];
- 2) **-Infinity($-\infty$)** (in **binary** form): [2] ;
- 3) **Smallest Negative Denormalized Value** (in **binary** form): [3];
- 4) **Smallest Positive Normalized Value** (in **binary** form): [4];

2. Convert the number $(-0.375)_{10}$ into the float12 representation (in **binary**). (3')

3. Assume we use IEEE **round-to-even** mode to do the approximation. Please calculate the addition: $(0\ 10000\ 010110)_2 + (0\ 10010\ 100100)_2$ and write the answer in **binary**. (The answer is represented in float12 as well) (3')

Problem 5: (23points)

<pre>char TX(char *str, int len) { int i = 0; char ret = 'A'; while (i < len) { switch (str[i]) { case 'a': ret += 1; break; case 'b': ret -= 32; break; case 'c': ret += 32; break; } } }</pre>	<pre>/* ASCII(0~9):0x30~0x39 * ASCII(A~Z):0x41~0x5a * ASCII(a~z):0x61~0x7a */ int main(void) { char str[4] = {'a','c','f','b'}; return TX((char *)str,4); }</pre>
---	--

<pre> case 'd': ret -= __[1]__; break; case __[2]__: ret = 3; __[3]__; default: ret = __[4]__; } i = i + 1; } return ret; } </pre>	<pre> .section .rodata .align 4 .L7: .long .L1 .long .L2 .long .L3 .long .L4 .long __[5]__ .long .L6 </pre>
<pre> <TX>: pushl %ebp movl %esp,%ebp pushl %ebx subl \$0x10,%esp __[6]_ \$0x41,%dl movl \$0x0, -0x8(%ebp) jmp .L10 .L0: addl __[7]_, -0x8(%ebp) .L10: movl -0x8(%ebp), %eax cmp 0xc(%ebp), %eax jl __[8]_ movzbl %dl, %eax addl \$0x10, %esp popl __[9]_ leave ret .L9: movl -0x8(%ebp), %ecx addl 0x8(%ebp), %ecx movzbl (%ecx), %ecx </pre>	<pre> subl __[10]_, %ecx cmp \$0x5, %ecx ja .L8 jmp __[11]_ .L1: addl \$0x1, %edx jmp .L0 .L2: __[12]_ \$0x20, %edx jmp .L0 .L3: __[13]_ \$0x20, %edx jmp .L0 .L4: subl \$0x4, %edx jmp .L0 .L6: __[14]_ \$0x3, %edx .L8: movl \$0x0, %ebx movb \$0x68, %bl movzbl %bl, %edx jmp .L0 </pre>

Suppose the C and assembly code are executed on a 32-bit little endian machine. Read the code and answer the following questions.

- Please fill in the blanks within C and assembly code. (1.5' * 14)
NOTE: no more than one instruction/statement per blank. If you think nothing is required to write, please write NONE.
- What is the return value of the **main** function? If it is a digit (0-9) or letter (A-Z, a-z), please use the digit or letter instead of its ASCII code. (2')

Problem 6: (27points)

One of the TA of ICS wrote a wrong Bubble Sort program and try to use gdb to find out the reason. The following C code and assembly code are executed on a **32-bit little endian** machine. In the C code, the definitions of array is omitted.

<pre>#include <stdio.h> int main(void) { int array[SIZE]; ... // initialize array sort(array, SIZE); return 0; }</pre>	<pre>void sort(int array[], int size) { int i; for (i=0; i<size-1; ++i) if (array[i] < array[i+1]) swap(&array[i], &array[i+1]); } void swap (int *left, int *right) { *left ^= *right; *right ^= *left; *left ^= *right; }</pre>
<pre>080483ed <swap>: 80483ed: 55 push %ebp 80483ee: 89 e5 mov %esp,%ebp 80483f0: 8b 45 08 mov 0x8(%ebp),%eax 80483f3: 8b 55 0c mov 0xc(%ebp),____[1]____ 80483f6: 8b 08 mov (%eax),%ecx 80483f8: 33 0a xor (%edx),%ecx 80483fa: 89 08 mov %ecx,(%eax) 80483fc: 33 0a xor (%edx),%ecx 80483fe: 89 0a mov %ecx,(%edx) 8048400: 31 08 xor %ecx,(%eax) 8048402: 5d pop %ebp 8048403: c3 ret 08048404 <sort>: 8048404: 55 push %ebp 8048405: 89 e5 mov %esp,%ebp 8048407: 57 push %edi 8048408: 56 push %esi 8048409: 53 push %ebx 804840a: 83 ec 08 sub \$0x8,%esp 804840d: 8b 7d 0c mov ____[2]____,%edi 8048410: 8d 47 ff lea -0x1(%edi),%eax 8048413: 85 c0 test %eax,%eax 8048415: 7e 2c jle 8048443 <sort+0x3f> 8048417: 8b 45 08 mov 0x8(%ebp),%eax</pre>	

804841a:	83 ef 01	sub \$0x1,%edi
804841d:	bb 00 00 00 00	mov \$0x0,%ebx
8048422:	8d 70 04	lea 0x4(%eax),%esi
8048425:	8b 50 04	mov 0x4(%eax),%edx
8048428:	39 10	cmp %edx,(%eax)
804842a:	7d 0c	jge 8048438 <sort+0x34>
804842c:	89 74 24 04	mov %esi,0x4(%esp)
8048430:	89 04 24	mov ____[3]____,(%esp)
8048433:	e8 b5 ff ff ff	call 80483ed <swap>
8048438:	83 c3 01	add \$0x1,%ebx
804843b:	39 fb	cmp %edi,%ebx
804843d:	74 04	je 8048443 <sort+0x3f>
804843f:	89 f0	mov %esi,%eax
8048441:	eb df	jmp 8048422 <sort+0x1e>
8048443:	83 c4 08	add \$0x8,%esp
8048446:	5b	pop %ebx
8048447:	5e	pop %esi
8048448:	5f	pop %edi
8048449:	5d	pop %ebp
804844a:	c3	ret
0804844b <main>:		
804844b:	55	push %ebp
804844c:	89 e5	mov %esp,%ebp
804844e:	83 ec 28	sub \$0x28,%esp
8048451:	c7 45 e0 09 00 00 00	movl \$0x9,-0x20(%ebp)
8048458:	c7 45 e4 0a 00 00 00	movl \$0xa,-0x1c(%ebp)
804845f:	c7 45 e8 0b 00 00 00	movl \$0xb,-0x18(%ebp)
8048466:	c7 45 ec 0c 00 00 00	movl \$0xc,-0x14(%ebp)
804846d:	c7 45 f0 0d 00 00 00	movl \$0xd,-0x10(%ebp)
8048474:	c7 45 f4 0e 00 00 00	movl \$0xe,-0xc(%ebp)
804847b:	c7 45 f8 0f 00 00 00	movl \$0xf,-0x8(%ebp)
8048482:	c7 45 fc ____[4]____	movl \$0x10,-0x4(%ebp)
8048489:	c7 44 24 04 08 00 00 00	movl \$0x8,0x4(%esp)
8048491:	8d 45 e0	lea -0x20(%ebp),%eax
8048494:	89 04 24	mov %eax,(%esp)
8048497:	e8 68 ff ff ff	call ____[5]____ <sort>
804849c:	b8 00 00 00 00	mov \$0x0,%eax
____[6]____:	c9	leave
80484a2:	c3	ret

1. Fill in the blanks in the Assembly Code. (1'*6=6').
2. What is the value of the size of array? (1')

3. **BEFORE** executing the instruction "push %ebp" (0x804844b), the value of register %esp is 0xffffd4fc. Please fill the following blanks.(1'*6=6')

AFTER executing the instruction "call __[5]__ <sort>" (0x8048497)

register	value
%esp	[1]
%ebp	[2]

AFTER executing the instruction "push %edi" (0x8048407)

register	value
%esp	[3]
%ebp	[4]

AFTER executing the instruction "ret" (0x804844a)

register	value
%esp	[5]
%ebp	[6]

4. We use **gdb** to run this program and step into the first loop of function **sort**. The **FIRST** time we stop before the execution of instruction at

```
8048428: 39 10      cmp %edx, (%eax)
804842a: 7d 0c      jge 8048438 <sort+0x34>
```

What are the values (in hexadecimal) of %edx and (%eax)? Which instruction will be executed after "jge 8048438"? (6')

5. After that, we continue the execution of **gdb**. After several loops in function **sort**, we stop again before the execution of instruction at

```
804842a: 7d 0c      jge 8048438 <sort+0x34>
```

We use **gdb** to print some registers. Suppose the start address of **array** is 0xffffd4d8. Please fill the table below. (2'*4=8')

register	value
%eax	0xffffd4e4
%ebx	__[1]__
%ecx	0x9
%edx	__[2]__
%esi	0xffffd4e8
%edi	0x7
%esp	0xffffd4b4
%ebp	__[3]__
%eip	__[4]__