

Homework 14

Replacement

Let's consider about a page-removal algorithm: clock algorithm. Suppose we have a primary device which has 3 physical blocks, every time a reference string P come, it will follow the pseudo-code:

```
if hit(P)
    res_block = block contains P
    referenced_bit[res_block] ← True
    clock_arm does not change
else
    while referenced_bit[clock_arm]
        referenced_bit[clock_arm] ← False
        clock_arm ← next block
    res_block = block[clock_arm]
    referenced_bit[clock_arm] ← True
    clock_arm ← next block
return res_block
```

Fill in the table (If you don't know what to fill just write down a '-'). Note: '*' means the position of the clock arm; you also need to tell what the referenced bit is for each page block at that time (1 for True and 0 for False) .

Time	0	1	2	3	4	5	6	7	8
Reference string	-	3	4	2	6	4	3	7	4
Primary Device Contents	*	3	3	3 *	6	6	6 *	6	4
	-	*	4	4	4 *	4 *	4	7	7 *
	-	-	*	2	2	2	3	3 *	3
Referenced Bit	0	1	1	1	1	1	1	0	1
	0	0	1	1	0	1	0	1	1
	0	0	0	1	0	0	1	1	0
Page Absent	-	Y	Y	Y	Y	N	Y	Y	Y

Locking

What are problems in the following simple implementation of lock?

```

typedef struct __lock_t {
    int flag;
} lock_t;
void init (lock_t *mutex) {
    mutex->flag = 0; // 0 -> lock is available 1 -> held
}
void lock (lock_t *mutex) {
    while (mutex->flag == 1) ; // spin-wait (do nothing)
    mutex->flag = 1; // now SET it!
}
void unlock(lock_t *mutex) { mutex->flag = 0; }

```

No mutual exclusion. Consider two threads doing lock, T1 may see mutex->flag is 0, then he is scheduled out. Then T2 sees mutex->flag is 0 too and set the flag to 1 and enter the critical section. Then T1 get scheduled in, he set the flag to 1 and enter the critical section too.

Scheduling

We have following jobs in the workload. No I/O issues are involved.

Job	Arrival Time	Run time
A	0ms	4ms
B	1ms	1ms
C	4ms	5ms
D	6ms	2ms

- ✧ When a job arrives, it is added to the tail of the work queue.
 - ✧ CPU picks job to run after all queue operations.
 - ✧ The **MLFQ** policy has 2 priority queues, higher one with time-slice of 1ms and lower one with time-slice of 2ms. We use **RR** in each queue. Priority boost isn't supported.
 - ✧ No preemption in **MLFQ**.
 - ✧ We do RR by moving the **recently executed task** to the end of the queue.
 - ✧ The priority of operations is RR movement > accepting new job.
- Please calculate the **average** response time and **average** turnaround time for different scheduling policies.

Scheduling Policy	Turnaround Time	Response time
FIFO	5ms	2ms

STCF	4ms	0.25ms
MLFQ	4.75ms	0ms